

SKILLS

Languages	Proficient: Python, SQL
Databases	Experience: SQL, NoSQL, GraphDB, VectorDB
Tech Stack	Git, Bash, Databricks, Jupyter Notebook (Anaconda), VS Code, Scikit-learn, PyTorch, FastAPI, Docker, Kubernetes
GenAI & LLMs	OpenAI API, LangChain, LangGraph, CrewAI, MCP, Prompt Engineering, RAG, Hybrid RAG, GraphRAG, Tool Calling, AI Agents, LangFuse/W&B
Cloud & MLOps	Azure, AWS, Google Cloud Platform (GCP), Azure OpenAI, Azure AI Foundry, Azure Container Apps, Azure Functions, CI/CD, Docker, Kubernetes
AI Governance	AI Guardrails, AI Governance, GRC, Responsible AI, Explainability, Auditability, Traceability, AI Evaluation, AI Observability, Human-in-the-Loop, Prompt Versioning, Policy Enforcement, Model Monitoring, Risk Assessment, ISO/IEC 42001, GDPR, EU AI Act
Tools & Process	Agile, Git, Visio, UAT, Test Data Creation, Stakeholder Management, Solution Design, REST APIs
Certifications	Microsoft Azure AI, Databricks, CrewAI, Hugging Face, DeepLearning.AI Prompt Engineering, LangGraph AI Agents, LLM Fine-tuning.

EXPERIENCE

AI Engineer – Network Intelligence & AI Platforms AWTG	Jul 2025 — Present London, UK
- Led the end-to-end delivery of enterprise AI solutions, transforming proof-of-concept ideas into production-ready applications by integrating LLMs, retrieval systems, workflow orchestration, secure APIs, and cloud-native deployment. - Scaled an AI-powered customer support platform for the British Council from an early proof of concept into a production solution supporting over 65,000 active users. Designed the retrieval architecture, workflow orchestration, AI guardrails, and evaluation pipelines, enabling grounded responses while minimizing manual customer support effort. - Collaborated on AI solutions supporting government border operations by developing conversational AI capabilities using LLMs, retrieval augmentation, workflow automation, and policy-aware decision support while ensuring governance, compliance, explainability, auditability, and end-to-end traceability throughout the decision-making process. - Addressed inconsistent AI responses across enterprise knowledge bases by designing Hybrid RAG and GraphRAG architectures utilizing knowledge graphs, vector search, metadata enrichment, and semantic retrieval, enhancing multi-hop reasoning, contextual retrieval, explainability, and evidence-backed answer generation. - Established enterprise AI governance by implementing AI guardrails, policy enforcement, prompt versioning, access controls, audit logging, traceability, and human-in-the-loop workflows, enabling secure, compliant, explainable, and production-ready AI systems. - Developed AI observability and evaluation frameworks that continuously monitored groundedness, hallucination rates, retrieval quality, latency, safety, operational cost, and user feedback while providing complete traceability between user prompts, retrieved evidence, tool execution, and generated responses. - Co-designed and developed KAI, an enterprise AI platform providing reusable agent orchestration, secure tool execution, memory management, retrieval services, workflow automation, and enterprise integrations, significantly accelerating AI solution delivery across multiple customer engagements. - Developed production-ready AI applications using GPT, Claude, Llama, and Mistral, integrating context engineering, structured outputs, tool calling, enterprise knowledge retrieval, and backend services to improve decision support and operational efficiency. - Successfully designed and evaluated an intelligent handover optimisation framework between Terrestrial Networks (TN) and Non-Terrestrial Networks (NTN) using Reinforcement Learning, deep learning, and Generative AI. Modelled network selection using mobility, radio, latency, throughput, and service continuity to optimise adaptive handover decisions beyond traditional rule-based approaches. - Integrated AI services with HubSpot, Salesforce, and enterprise APIs to automate CRM synchronisation, lead qualification, customer engagement workflows, and email-triggered actions while maintaining secure authentication, governance, and auditability. - Applied modern software engineering practices including Python, PyTorch, Git, Docker, Kubernetes, CI/CD, Azure Container Apps, Azure Functions, and automated testing to deliver scalable, secure, and production-ready AI systems while mentoring engineers on GraphRAG, enterprise retrieval, AI evaluation, governance, and production deployment.	
Senior AI/ML Engineer Accenture	May 2020 — Sept 2024 Pune, India
Delivered end-to-end AI and machine learning solutions for enterprise clients, translating complex business problems into scalable production systems supporting forecasting, optimisation, conversational intelligence, and decision support.	

+44 75 0238 7536
London, UK
kandelsudip245@gmail.com

Sudip Kandel

Senior AI/ML Engineer

Portfolio: [Click here](#)
Recommendations: [Click here](#)
github.com/sudipkandel123
linkedin.com/in/sudipkandel123

- Built conversational AI and enterprise knowledge assistants using **LangChain**, Hugging Face Transformers, retrieval-based architectures, and **Python**, improving response quality and enabling natural-language access to enterprise information.
- Designed scalable data pipelines for ingestion, validation, transformation, and feature engineering across structured and unstructured datasets, improving data quality and reducing manual processing effort.
- Developed machine learning solutions for forecasting, anomaly detection, optimisation, predictive analytics, and business intelligence, enabling data-driven operational and strategic decision-making.
- Implemented automated model evaluation, regression testing, quality assurance, and monitoring pipelines that improved deployment confidence and reduced production defects.
- Built reusable backend APIs and inference services that integrated machine learning models into enterprise applications and operational workflows.
- Partnered with product owners, business analysts, and technical stakeholders to define success metrics, prioritise delivery roadmaps, and communicate AI solutions to both technical and executive audiences.
- Applied software engineering best practices including Git, automated testing, documentation, and CI/CD pipelines to deliver maintainable, scalable, and production-ready AI applications.

EDUCATION

MSc. Artificial Intelligence & Machine Learning , <i>Liverpool John Moores University</i>	Jun 2026
MSc. Management of Business Information Technology (83%) , <i>University of Greenwich</i>	Sept 2025
Post Graduate Diploma. Artificial Intelligence and Machine Learning , <i>IIIT Bangalore (IIIT-B)</i>	Feb 2025
Bachelor of Engineering. Information Science, First Class Honours (91%) , <i>New Horizon College of Engineering (NHCE)</i>	May 2019

ACTIVITIES

- **LangChain Ambassador (UK)** – Organise AI meetups, workshops, and hackathons while delivering talks on Agentic AI, GraphRAG, AI governance, and production AI to engineers and industry professionals.
- Mentored and educated **hundreds of engineers, students, and professionals** through AI workshops, hackathons, and technical talks, helping teams adopt AI responsibly and transition solutions from proof of concept to production.
- Active AI writer on [Medium](#) and [LinkedIn](#). Winner of multiple [AI Hackathons](#) focused on solving real-world problems.
- Recipient of multiple recognitions including [GenAI Implementation Award](#), [Accenture Outstanding Employee Award](#), [EXL Distinguished Employee Award](#), and [NHCE Techorizon](#).
- Published research on [Election Sentiment Analysis](#) (2019), applying NLP and machine learning techniques to social media analysis.
- Passionate about building trustworthy AI that delivers measurable business and societal impact through responsible innovation, governance, and human-centred design.
- Outside work, volunteer as a [Wellness Coach](#), [Meditation Instructor](#), [Bamboo Flautist](#), and community [Volunteer](#), promoting wellbeing and supporting underserved communities.